

Pricing the Print: Do Prediction Markets Tell You When to Sell CPI-Day Volatility?

Quant research note · October 2026 · Code and data: <https://github.com/shmoney2/kalshi-cpi-prereg>

Summary

Selling S&P 500 options into a CPI release collects an event premium, but a few releases produce moves large enough to erase months of gains. We tested whether Kalshi prediction markets tell a seller when to stand aside. Both tests were pre-registered and timestamped on Solana mainnet. **Stage 1** (49 releases, 2022-07 to 2026-09): Kalshi's uncertainty about the next FOMC decision does not reliably predict how hard stocks react to CPI surprises (**NO-GO**). **Stage 2-lite** (50 releases, VIX1D as the option price): the realized move was smaller than implied on 74% of releases, yet the average variance premium is about zero, a defined-risk butterfly is flat in sample, and neither Kalshi filter improved timing. The CPI-day premium looks like fair compensation for tail risk; we would not trade it, and the next step is real option quotes. What is distinctive is the design: Kalshi's CPI and Fed ladders map onto a news-times-sensitivity view of event premia, both plans were fixed and timestamped on a public blockchain before the outcome data, and every number reproduces from one command.

Hypothesis

An announcement's market impact is the news it releases times the market's sensitivity to that news (Knox, Londono, Samadi and Vissing-Jorgensen, "Equity Premium Events"). Kalshi's CPI ladders price the news; its Fed ladders say whether the next FOMC meeting is "live", which should raise sensitivity. **H1:** in $r = a + b \cdot S + c \cdot (S \times L) + d \cdot L + e \cdot VIX$, where r is the S&P 500 move from the prior close to 09:35 ET, S the core CPI surprise against Kalshi's 08:25 expectation and L Fed liveness the evening before, $c < 0$: a hot print hurts more when the meeting is live. **H2 (mechanism):** CPI moves Kalshi's expected Fed rate more when the meeting is live. **Stage 2 hypothesis:** skipping releases where a Kalshi input is in the top third of its history avoids the releases where selling loses, better than random skips and better than filtering on the option price.

Who is on the other side. Buyers of CPI-day options are largely hedgers paying for protection against a jump on the print. A premium above expected losses could persist because few investors are willing to hold jump risk, so those who do are paid for it. The trade fails if the premium only pays for the tail risk itself: sellers then break even on average once the rare large moves arrive, and no timing filter can rescue it. Stage 2 tests exactly this failure condition. A Kalshi filter could add to the option price because VIX1D prices how large a move to expect but not why: Kalshi's CPI ladder gives the distribution of the print, and its Fed ladder shows whether a surprise would change the next policy decision, the channel that should make some releases riskier than their implied volatility suggests.

Pre-registration. Stage 1's plan and code were frozen in commit 1c084b5 and stamped on Solana mainnet at 2026-10-03 17:51 UTC ([transaction 5WzPtpFrFd...](#)) before any release-day data were downloaded. Stage 2-lite was frozen in 02e251b and stamped at 2026-10-03 20:51 UTC ([transaction 2sDcsGVpQp...](#)) before VIX1D was downloaded, but after Stage 1's returns had been computed; its plan says so. Each memo holds the SHA-256 of a manifest hashing every frozen file, and `timestamp_solana.py` verify rechecks it.

Data

- **Kalshi public API:** core and headline CPI ladders (KXCPI, KXCPICORE) and next-meeting Fed ladders (KXFED), 2021–2026, snapshots at 15:45 ET the evening before (t_0), 08:25 and 10:00 on release day. Price rule: last trade within 24 hours, else the bid-ask midpoint if the spread is at most \$0.10; ladders need at least 4 usable strikes.
- **Massive (formerly Polygon.io):** SPY daily and minute bars and dividends. SPY replaces the index, which the data plan lacked; the switch was fixed before the freeze. Dividends are added back on ex-dates.
- **Cboe:** VIX and VIX1D daily history (free files). From 16:00 to 16:15 VIX1D uses only next-day SPX options, so its close on a release eve prices release-day risk; values before its April 2023 launch are Cboe back-calculations, starting May 2022.
- **BLS CPI release calendar:** September 2025 CPI (moved by the shutdown after Kalshi's markets had closed) and October 2025 CPI (never published) are excluded.

Samples: 49 releases with complete Stage 1 data (2022-07-13 to 2026-09-11); 50 with VIX1D for Stage 2 (from 2022-06-10). Three data-plumbing fixes made after the freeze (matching events to the BLS calendar, a standing-quote lookback, reading a non-numeric settlement value) were committed with tests before the test was run; Appendix E lists them with two parsing fixes made before the freeze.

Methodology

Stage 1. H_1 is estimated by OLS with HC3 standard errors and one-sided tests; L is the standard deviation of Kalshi's distribution for the next meeting's upper bound, standardized; VIX is the prior close. GO requires all five of: $c < 0$; $HC3\ p < 0.10$; within-year permutation $p < 0.10$ (10,000 draws); the same sign in at least 90% of leave-one-out fits; and $c < 0$ without 2022. Eight robustness variants and H_2 are reported but never decide.

Stage 2-lite. The implied daily move is $\sigma = VIX1D/\sqrt{252}$ and the variance premium $\sigma^2 - r^2$, with r the close-to-close SPY return; positive means a variance seller earned money. The tradeable proxy is a short at-the-money butterfly with wings two implied moves out, priced from a normal distribution with standard deviation σ , with entry cost at 3% of the straddle (0% and 6% also run) and 1% of capital at risk per release. The 3% cost is an assumption: at the median release ($\sigma = 1.10\%$, straddle 0.88% of the index, $SPX \approx 5,395$) it equals 2.6 bps of index notional, about 1.4 index points for the position or 0.35 points per leg across four legs. A filter skips a release when its input is above the two-thirds point of all earlier values, after a 6-release burn-in, so every decision uses only the past. "Kalshi improved the decision" needs all three: skipped releases earn a lower premium than kept ones; a random skip set of the same size does as well in under 10% of 10,000 draws; and the filter beats one built on VIX1D alone. After Stage 1's NO-GO the primary input is Kalshi's core CPI uncertainty (ladder standard deviation); Fed liveness, CPI uncertainty measured by interquartile range, news $\times$ sensitivity, VIX1D and VIX filters are reported.

Track additions (post hoc, logged as dated deviations). Out of sample is the most recent 20% of each sample's date span. Annualized figures assume 12 releases a year and count skipped releases as zero; turnover is option premium traded per year as a multiple of capital. Full-sample results were seen before the split, which follows the track's mechanical rule. A pricing sensitivity reprices the butterfly under Student-t distributions scaled to VIX1D's variance.

Multiple testing. Every test, variant, strategy and cost level run is counted in Table 1 (itemized in Appendix D). Only the Stage 1 primary model and the Stage 2 primary filter can produce a positive verdict.

Group	Specifications	Can decide
Stage 1, pre-registered	17	6 (primary model, 5 conditions)
Stage 1, data diagnostics	2	none
Stage 2, pre-registered	14	3 checks on the primary filter
Post hoc (split, pricing sensitivity)	43	none
Total	76	

Table 1. Specifications run on real data. Resampling draws count once per test.

Results

Sample	n	c	HC3 SE	p (HC3)	p (perm.)
Full sample (decision)	49	-1.62	1.61	0.159	0.754
In sample, to 2025-11-10	39	-1.23	1.79	0.249	0.710
Out of sample (descriptive)	10	1.52	5.19	0.609	0.647

Table 2. Stage 1 primary model; c is per 1 SD of liveness, % per percentage point of surprise. About 10 releases cannot support inference.

Stage 1: NO-GO. The interaction has the predicted sign: a core print 0.1 point hotter than Kalshi expected moved SPY -0.19% when the next meeting looked settled and -0.52% when it looked live. But it fails three of five conditions (Table 2): the HC3 and permutation p-values are 0.16 and 0.75, and without 2022 the sign reverses ($c = 4.39$), so the effect rests on 2022's hiking cycle. The mechanism does hold: CPI moves Kalshi's expected Fed rate more when the meeting is live (H2, $c = 6.71$, $p = 0.007$), partly by construction since a wider distribution has more room to move. The market's odds react; stock sensitivity does not reliably follow.

Stage 2-lite: no evidence that Kalshi helps. Realized variance was below VIX1D-implied variance on 74% of releases (the average move was 0.70 implied standard deviations), yet the mean variance premium is -0.20 (90% interval -1.36 to 0.80): 2 releases with moves above two implied standard deviations offset the many small wins. After costs and wing premium the butterfly made money on 60% of releases. The CPI-uncertainty filter skipped releases that were worse for sellers and kept better ones than the VIX1D filter, but random skips did as well 17% of the time ($p = 0.174$), so it fails the decision rule. Always selling was flat in sample (0.00% a year at 3% cost) and positive only over the 10 out-of-sample releases (4.29% a year; Table 3, Figure 1). With 10 monthly observations an annualized Sharpe ratio has a standard error of about $\sqrt{(12/10)} \approx 1.1$ (1.3 at the observed 2.85, by Lo's formula), so the out-of-sample Sharpe ratios rest on very little data. This is the Hypothesis's failure condition: the premium appears to pay for the tails rather than exceed them.




Figure 1. Growth of capital at 3% entry cost; the shaded region is out of sample.

Cost	Strategy	Per.	Trades	Ann. ret.	Ann. vol.	Sharpe	Max DD	Turnover
3%	Always sell	IS	40	0.00%	2.02%	0.00	3.25%	0.12×
3%	Always sell	OOS	10	4.29%	1.50%	2.85	0.57%	0.12×
3%	Kalshi filter (CPI SD)	IS	36	0.09%	1.92%	0.05	3.25%	0.11×
3%	Kalshi filter (CPI SD)	OOS	9	3.57%	1.52%	2.35	0.57%	0.11×
3%	Option price (VIX1D)	IS	36	0.11%	1.89%	0.06	3.25%	0.11×
3%	Option price (VIX1D)	OOS	7	3.17%	1.04%	3.05	0.05%	0.08×
6%	Always sell	IS	40	−0.32%	1.96%	−0.16	3.32%	0.11×
6%	Always sell	OOS	10	3.85%	1.46%	2.63	0.58%	0.11×
6%	Kalshi filter (CPI SD)	IS	36	−0.20%	1.87%	−0.11	3.32%	0.10×
6%	Kalshi filter (CPI SD)	OOS	9	3.19%	1.47%	2.17	0.58%	0.10×
6%	Option price (VIX1D)	IS	36	−0.18%	1.84%	−0.10	3.32%	0.10×
6%	Option price (VIX1D)	OOS	7	2.86%	0.98%	2.90	0.08%	0.08×

Table 3. In sample (IS, 40 releases) and out of sample (OOS, 10 releases from 2025-12-18), 1% of capital at risk per release. All seven strategies are in Appendix A.

Risk

The payoff is short tail risk with a hard cap: the worst release loses the full 1% at risk, which happened whenever the move passed a wing. All strategies share the same 3.25% in-sample drawdown from the burn-in releases of mid-2022. The largest risk is model risk. The normal model is the pre-registered proxy; Table 4 stress-tests tail pricing with Student-t models, whose fatter tails make wings dearer. Real SPXW quotes put the wings much closer to the normal model, with straddles and wings both below it, and crossing cost 0.7% of the straddle, not 3%. Either way the result is flat in sample.

Butterfly pricing	3%:	6%:
	mean/trade	mean/trade
Normal (pre-registered)	+0.072%	+0.043%
Student-t, 6 df	−0.003%	−0.028%
Student-t, 4 df	−0.060%	−0.084%

Table 4. Always sell, all 50 releases; post-hoc sensitivity, reported only.

Factor exposure. Regressing always-sell's per-release return (% of capital) on SPY's release-day return, its absolute value, the release-day VIX change and VIX1D on the eve (HC3 errors, $R^2 = 0.77$) gives a market beta of 0.03 ($t = 0.2$) and −0.55 per 1% absolute move ($t = -3.3$); the VIX change adds nothing ($t = 0.5$), and returns rise with the VIX1D level ($t = 3.5$). Raw correlations with SPY's return and its absolute value are 0.02 and −0.80: the position is short the size of the move, not the market.

Regimes and stress. The 2022 hiking cycle was the bad regime (Table 5). The three worst releases, 2022-09-13 (SPY −4.45%), 2022-11-10 (SPY +5.35%), 2023-11-14 (SPY +1.92%), each lost the full 1%, and any move past the wings, whether 3 or 5 implied standard deviations, loses exactly 1%, so a crash loss is bounded by construction. Maximum losses hit 3 of 50 releases and never twice in a row (longest losing run 4); 7 maximum losses in a row would trigger the halve-size rule and 14 the stop.

Regime (always sell, 3% cost)	n	Mean/trade	Total	Worst	Hit rate
2022	7	−0.392%	−2.72%	−1.00%	29%
2023–2026	43	+0.147%	+6.46%	−1.00%	65%
VIX above 17.7 (median)	25	+0.086%	+2.12%	−1.00%	60%
VIX at or below 17.7	25	+0.057%	+1.41%	−1.00%	60%

Table 5. Always sell by regime; VIX regimes split at the median prior close.

Exposure limits and operating rules (set after the backtest, not pre-registered): one position at a time, since each butterfly expires the day after it is opened; at most 1% of capital lost per release and 12% at risk in a year; skip any release whose date moves after the position is planned, as September 2025 did; halve position size after a 6% drawdown and stop trading at 10%. None would have triggered: the worst drawdown of any strategy was 3.25% at 3% cost (3.32% at 6%).

Liquidity and capacity

The Kalshi inputs are read, not traded, so Kalshi liquidity limits signal quality rather than capacity. At t_0 the median 24-hour volume was about 7,204 contracts across the core CPI ladder and 11,910 across the Fed ladder, with median quoted spreads of \$0.04 and \$0.01; thin far strikes still shape the ladder's standard deviation (see Limitations). On the options side, the strategy trades SPX options expiring the next day, about once a month. At 1% of capital at risk per release it trades option premium worth roughly $0.12\times$ capital a year, so capacity is unlikely to bind before the edge does. At the median release one SPX butterfly (\$100 multiplier) can lose about 51.5 index points, or \$5,152, so risking 1% per release implies about \$515,000 of capital per contract; XSP, one-tenth the size, brings this to \$515 and about \$51,500. Moving the entry-cost assumption from 3% to 6% lowers always-sell's mean return per trade from 0.072% to 0.043%.

Measured SPXW costs. Measured after the main results were known, under a separate plan stamped on Solana before any release-eve option quotes were downloaded ([transaction](#)), we read real SPXW quotes and volume from 15:45 to 16:15 on each of the 50 release eves. It computes no returns. All four legs had usable quotes on 48 eves; on the other 2 the call-wing strike was never quoted, and missing quotes were not filled in. Crossing all four spreads cost a median 0.7% of the straddle (90th percentile 0.9%, highest 1.2%), or 0.32 index points against 1.34 under the 3% assumption, so 3% was conservative and 6% is a remote stress. Depth is the tighter limit: the thinnest top-of-book size was a median 11 contracts, so larger orders would walk the book. Trading at most 10% of the thinnest leg's eve volume (median 1,897 contracts) allows a median 190 butterflies, about \$87 million of capital at 1% risk, and 6 on the thinnest eve. At 16:15 real straddle mids were a median 0.96 of the normal-model price from VIX1D (IQR 0.92–0.99) and wing mids 0.84 (0.72–0.94).

Limitations and what didn't work

About 50 releases give low power: "NO-GO" and "no evidence" do not prove the effects are absent, and Stage 1's pre-registered power table shows effects smaller than a doubling of sensitivity are often missed. The ladder standard deviation used for liveness and CPI uncertainty partly measures how widely Kalshi lists strikes (correlation 0.76 with ladder span for liveness). The Stage 2 returns are proxies from VIX1D and a pricing model, the close-to-close window includes the whole release day, and pre-2023 VIX1D values are back-calculated. The out-of-sample split came after the full-sample results.

- Fed liveness as a gauge of stock sensitivity to CPI (Stage 1 NO-GO; the sign reverses without 2022).
- A Kalshi CPI-uncertainty filter (random-skip $p = 0.17$); Fed liveness and IQR variants did no better than always selling.
- A reliable always-sell edge: flat in sample; a loss only under the Student-t stress test (see Risk).
- Index and futures data: the data plan had no SPX index data, and a futures reaction window was dropped because minute data could not be exported with verifiable timestamps.

Conclusion. In 2022–2026 the CPI-day premium in S&P 500 options looks like fair compensation for tail risk, and Kalshi's odds do not identify the releases to avoid. We would not trade this strategy. The next step is Stage 2 with real SPXW prices at entry and settlement, stamped before those data are downloaded.

References

- Bureau of Labor Statistics. Consumer Price Index release schedule and news release archive.
https://www.bls.gov/schedule/news_release/cpi.htm; <https://www.bls.gov/bls/news-release/cpi.htm>
- Bürgi, C., Deng, W. and Whelan, K. (2026). "Makers and Takers: The Economics of the Kalshi Prediction Market." CEPR Discussion Paper 20631; CESifo Working Paper 12122.
- Cboe Global Markets. Cboe Volatility Index Mathematics Methodology.
https://cdn.cboe.com/resources/indices/Cboe_Volatility_Index_Mathematics_Methodology.pdf
- Cboe Global Markets. "What the VIX and VIX1D Indices Attempt to Measure and How They Differ" (VIX1D calculation, including the 4:00 to 4:15 p.m. next-day-only period). <https://www.cboe.com/insights/posts/what-the-vix-and-vix-1-d-indices-attempt-to-measure-and-how-they-differ>
- Cboe Global Markets. VIX historical data, VIX_History.csv.
https://cdn.cboe.com/api/global/us_indices/daily_prices/VIX_History.csv
- Cboe Global Markets. VIX1D historical data, VIX1D_History.csv.
https://cdn.cboe.com/api/global/us_indices/daily_prices/VIX1D_History.csv
- Diercks, A. M., Katz, J. D. and Wright, J. H. (2026). "Kalshi and the Rise of Macro Markets." Finance and Economics Discussion Series 2026-010, Board of Governors of the Federal Reserve System; NBER Working Paper 34702.
- Kalshi. Trade API v2, public market data (series KXCPI, KXCPCICORE, KXFED).
<https://api.elections.kalshi.com/trade-api/v2>
- Knox, B., Londono, J. M., Samadi, M. and Vissing-Jorgensen, A. "Equity Premium Events." Federal Reserve Board working paper.
- Lo, A. W. (2002). "The Statistics of Sharpe Ratios." Financial Analysts Journal 58(4), 36–52.
- Massive (formerly Polygon.io). Stocks aggregates and reference dividends for SPY.
<https://api.massive.com>
- Solana mainnet memo transactions of the two pre-registrations (see Hypothesis) and the liquidity measurement plan (see Liquidity and capacity); verifiable with `timestamp_solana.py` verify.
- Project repository, code and derived data: <https://github.com/shmoney2/kalshi-cpi-prereg>

Appendix

A1. All Stage 2 strategies, in and out of sample, 3% entry cost

Strategy	Per.	Trades	Ann. ret.	Ann. vol.	Sharpe	Max DD	Turnover
Always sell	IS	40	0.00%	2.02%	0.00	3.25%	0.12×
Always sell	OOS	10	4.29%	1.50%	2.85	0.57%	0.12×
Kalshi filter (CPI SD)	IS	36	0.09%	1.92%	0.05	3.25%	0.11×
Kalshi filter (CPI SD)	OOS	9	3.57%	1.52%	2.35	0.57%	0.11×
News × sensitivity	IS	40	0.00%	2.02%	0.00	3.25%	0.12×
News × sensitivity	OOS	9	3.57%	1.52%	2.35	0.57%	0.11×
CPI IQR (rep.)	IS	39	-0.06%	2.01%	-0.03	3.25%	0.11×
CPI IQR (rep.)	OOS	8	2.81%	1.47%	1.92	0.57%	0.09×
Fed liveness (rep.)	IS	37	-0.15%	1.90%	-0.08	3.25%	0.11×
Fed liveness (rep.)	OOS	8	3.71%	1.55%	2.39	0.57%	0.09×
Option price (VIX1D)	IS	36	0.11%	1.89%	0.06	3.25%	0.11×
Option price (VIX1D)	OOS	7	3.17%	1.04%	3.05	0.05%	0.08×
VIX only	IS	36	-0.12%	1.91%	-0.06	3.25%	0.11×
VIX only	OOS	7	3.17%	1.04%	3.05	0.05%	0.08×

A2. All Stage 2 strategies, in and out of sample, 6% entry cost

Strategy	Per.	Trades	Ann. ret.	Ann. vol.	Sharpe	Max DD	Turnover
Always sell	IS	40	-0.32%	1.96%	-0.16	3.32%	0.11×
Always sell	OOS	10	3.85%	1.46%	2.63	0.58%	0.11×
Kalshi filter (CPI SD)	IS	36	-0.20%	1.87%	-0.11	3.32%	0.10×
Kalshi filter (CPI SD)	OOS	9	3.19%	1.47%	2.17	0.58%	0.10×
News × sensitivity	IS	40	-0.32%	1.96%	-0.16	3.32%	0.11×
News × sensitivity	OOS	9	3.19%	1.47%	2.17	0.58%	0.10×
CPI IQR (rep.)	IS	39	-0.37%	1.96%	-0.19	3.32%	0.11×
CPI IQR (rep.)	OOS	8	2.48%	1.42%	1.75	0.58%	0.09×
Fed liveness (rep.)	IS	37	-0.45%	1.85%	-0.24	3.32%	0.11×
Fed liveness (rep.)	OOS	8	3.36%	1.50%	2.24	0.58%	0.09×
Option price (VIX1D)	IS	36	-0.18%	1.84%	-0.10	3.32%	0.10×
Option price (VIX1D)	OOS	7	2.86%	0.98%	2.90	0.08%	0.08×
VIX only	IS	36	-0.41%	1.86%	-0.22	3.32%	0.10×
VIX only	OOS	7	2.86%	0.98%	2.90	0.08%	0.08×

"rep." marks reported-only variants. Annualized at 12 releases a year; skipped releases count as zero.

B. Stage 1 decision conditions, robustness, mechanism and placebo

Condition	Passed	Detail
Hot prints hurt more when the meeting is live ($c < 0$)	yes	$c = -1.624$
HC3 one-sided p below 0.1	no	$p = 0.159$
Permutation one-sided p below 0.1	no	$p = 0.754$
Same sign after dropping any one release (at least 90%)	yes	100% of fits
Same sign without 2022	no	$c = 4.388$

Variant	n	c	SE	p
Headline instead of core surprise	52	-1.47	1.13	0.101
Close-to-close return (FOMC-day releases dropped)	48	-1.64	3.15	0.302
Sign-free: return on surprise (expects $c > 0$)	49	2.45	2.78	0.191
Liveness = 1 - modal probability	49	0.84	2.15	0.651
Liveness = entropy	49	1.21	2.09	0.716
No VIX control	49	-2.05	1.57	0.099
Surprise measured at t0 instead of 08:25	49	-1.40	1.72	0.209
Return scaled by VIX-implied daily move	49	-0.65	1.09	0.276

Mechanism (H2): $c = 6.71$ (SE 2.64), one-sided $p = 0.007$, $n = 49$. Placebo: ordinary-day volatility on liveness, coefficient -0.075 (SE 0.049), $n = 52$. 90% wild-bootstrap interval for c : -3.58 to 0.29 .

C. Stage 2-lite full-sample strategies and decision checks (3% cost)

Strategy	Trades	Mean var.		Total	Max DD
		prem.	Mean ret./trade		
Always sell	50	-0.200	0.07%	3.6%	2.9%
Kalshi filter (CPI SD)	45	-0.223	0.07%	3.3%	2.9%
News $\times$ sensitivity	49	-0.214	0.06%	2.9%	2.9%
CPI IQR (rep.)	47	-0.279	0.05%	2.1%	2.9%
Fed liveness (rep.)	45	-0.213	0.06%	2.5%	2.9%
Option price (VIX1D)	43	-0.303	0.07%	3.0%	2.9%
VIX only	43	-0.232	0.05%	2.2%	2.9%

Check	Passed	Detail
Releases Kalshi skips are worse for sellers than the ones it keeps	yes	skipped 0.005 vs kept 0.763
Better than skipping the same number at random ($p < 0.1$)	no	$p = 0.174$
Better than filtering on the option price alone	yes	-0.223 vs -0.303

Cost	Strategy	Mean return per trade
0%	Always sell	0.10%
0%	Kalshi filter (CPI SD)	0.10%
3%	Always sell	0.07%
3%	Kalshi filter (CPI SD)	0.07%
6%	Always sell	0.04%
6%	Kalshi filter (CPI SD)	0.04%

D. Every test, variant, strategy and cost level run

Stage	Category	Item	Specs	Decides
Stage 1	Pre-registered	Primary model H1 (decides GO/NO-GO)	1	yes
Stage 1	Pre-registered	H1 decision conditions: sign, HC3 p, within-year permutation p (10,000), leave-one-out sign share, fit excluding 2022	5	yes
Stage 1	Pre-registered	90% wild-bootstrap interval for c (9,999 draws)	1	no
Stage 1	Pre-registered	Robustness variants (headline, close-to-close, sign-free H3, 1 - modal, entropy, no VIX, surprise at t0, VIX-scaled return)	8	no
Stage 1	Pre-registered	Mechanism H2 (Kalshi expected rate response)	1	no
Stage 1	Pre-registered	Placebo (ordinary-day volatility on liveness)	1	no
Stage 1	Pre-registered	Consensus gap H4 (not run: no consensus data)	0	no
Stage 1	Data diagnostic	Liveness from 7-day vs fixed quote lookback (correlation, size of changes)	1	no
Stage 1	Data diagnostic	Ladder width vs liveness SD (correlation with strike span, tick-quoted strikes)	1	no
Stage 2	Pre-registered	Strategies: always sell, Kalshi filter (primary), news x sensitivity, CPI uncertainty IQR, Fed liveness, option price only, VIX only	7	primary only
Stage 2	Pre-registered	Decision checks (skipped worse than kept; random-skip permutation, 10,000; beats option-price filter)	3	yes
Stage 2	Pre-registered	90% bootstrap interval for the mean variance premium (5,000 draws)	1	no
Stage 2	Pre-registered	Entry-cost levels (0%, 3%, 6%) for always sell and the Kalshi filter	3	no
Post hoc	Track rule	Stage 1 primary model, in-sample and out-of-sample fits (each with permutation p)	2	no
Post hoc	Track rule	Stage 2 strategy x period x cost cells (7 x 2 x 2), six metrics each	28	no
Post hoc	Sensitivity	Butterfly pricing under Student-t (6 and 4 df) at 3% and 6% costs, always sell	4	no
Post hoc	Risk	Factor exposure of always-sell returns (four-factor regression and SPY beta)	2	no
Post hoc	Risk	Regime results (2022, 2023 to 2026, VIX above and below median)	4	no
Post hoc	Risk	Stress scenarios (worst releases, 3 to 5 SD shocks, loss runs against the size rules)	3	no
Total			76	

E. Disclosures, deviations and data-plumbing fixes

Disclosed before the Stage 1 freeze (PREREGISTRATION.md):

- The author viewed the 2022-06-10 CPI reaction in a thinkorswim replay while checking futures data; that release is outside the primary sample. Access checks requested SPY bars for two CPI days and printed only row counts.

Parsing fixes made before the freeze, part of the frozen code (commit 1c084b5):

- Strikes read from ticker suffixes for pre-2023 Kalshi contracts without strike fields; CPI events that closed the evening before a release dated to the release morning.

Fixed after the freeze, each committed with a test before the test was run:

- Kalshi events matched to the official BLS calendar; the two shutdown releases excluded (commit fd9e61c).
- Standing Kalshi quotes found with a 120-day daily-candle lookback, since Kalshi emits candles only for active periods (commit a3e3199).
- A settled value read from contract results when Kalshi's settlement field was not numeric (commit a3e3199).

Logged deviations (2026-10-03, after full-sample results were seen; reporting only, no decision rule changed): the out-of-sample split, the Student-t pricing sensitivity, the specification count and the reproduction script. Stage 2-lite added two reported-only filters before its stamp. The liquidity measurement (planned 2026-10-03, stamped 2026-10-04 before any release-eve quotes were downloaded) is a logged exception to the rule that options data wait for a Stage 1 GO; it computes no returns.

F. Dashboards

The interactive Stage 1 and Stage 2-lite dashboards are in the repository as dashboards/report.html and dashboards/stage2_report.html (<https://github.com/shmoney2/kalshi-cpi-prereg/tree/main/dashboards>). Download them and open them in a browser.